

Foundations of Modern Data Systems (FOMO)

Winter Semester 2025/2026
Exercise Session #2

23/10/25

Outline

1. Hands-on session: **Experiments & Profiling (con't)**
2. The memory hierarchy – practical view
3. Project 1
4. Hands-on session: **Project 1 code walk-through**

Quick PSA

- The moodle form to provide us with your ssh-key closes the day of the course registration deadline, don't be late!

Hands-on session

Experiments & Profiling (con't)

A bit more about PMUs & profiling

- Using some slides from our friends at TU Dortmund, presented at BTW2025

<https://tinyurl.com/nodmc-prof>



Tutorial: Understanding Application Performance on Modern Hardware

Profiling Foundations and Advanced Techniques

Jan Mühlig, Roland Kühn, and Jens Teubner

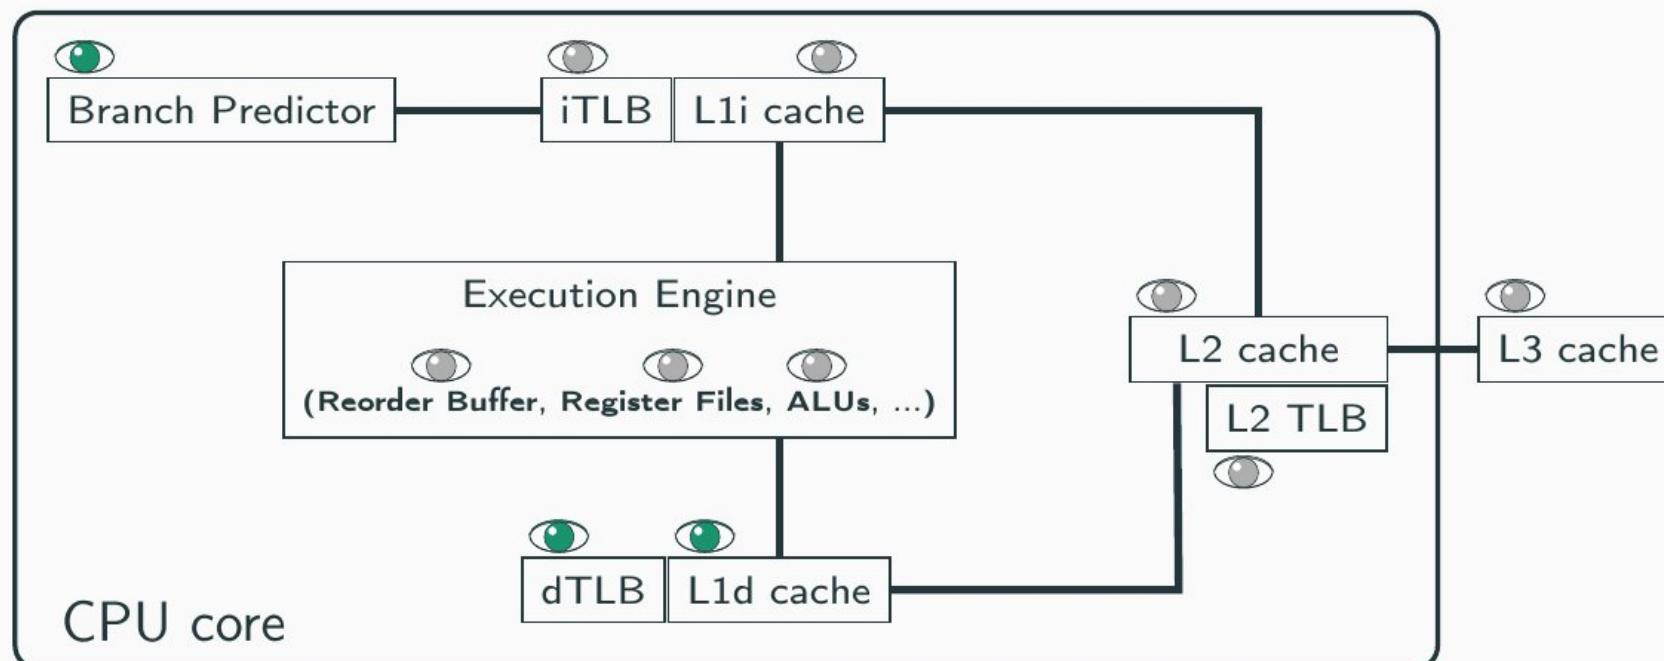
March 04, 2025

DBIS Group, TU Dortmund University

Performance Monitoring Units



- Modern CPUs are equipped with *Performance Monitoring Units*
- (Most) PMUs can **track** multiple events simultaneously
 - E.g., L1d hits/misses, dTLB hits/misses, branch predictions, ...
 - Hundreds to thousands events available (depending on the CPU vendor and model)



Profiling: Statistics



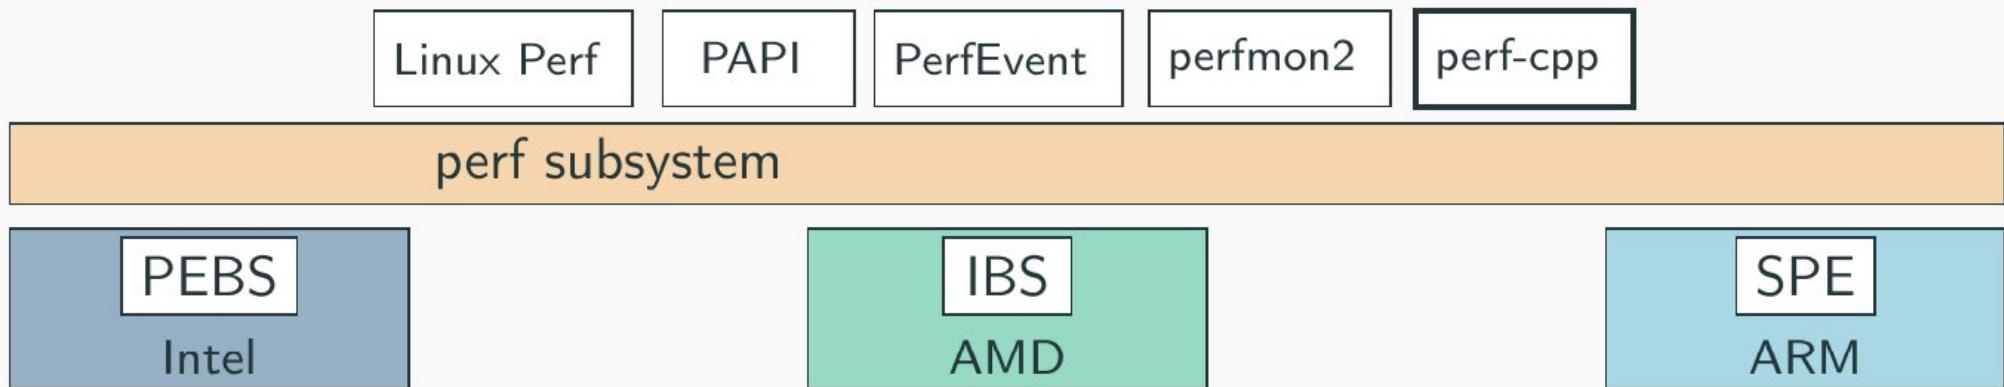
```
> perf stat ./my-app
```

read
performance
counter

read
performance
counter
again



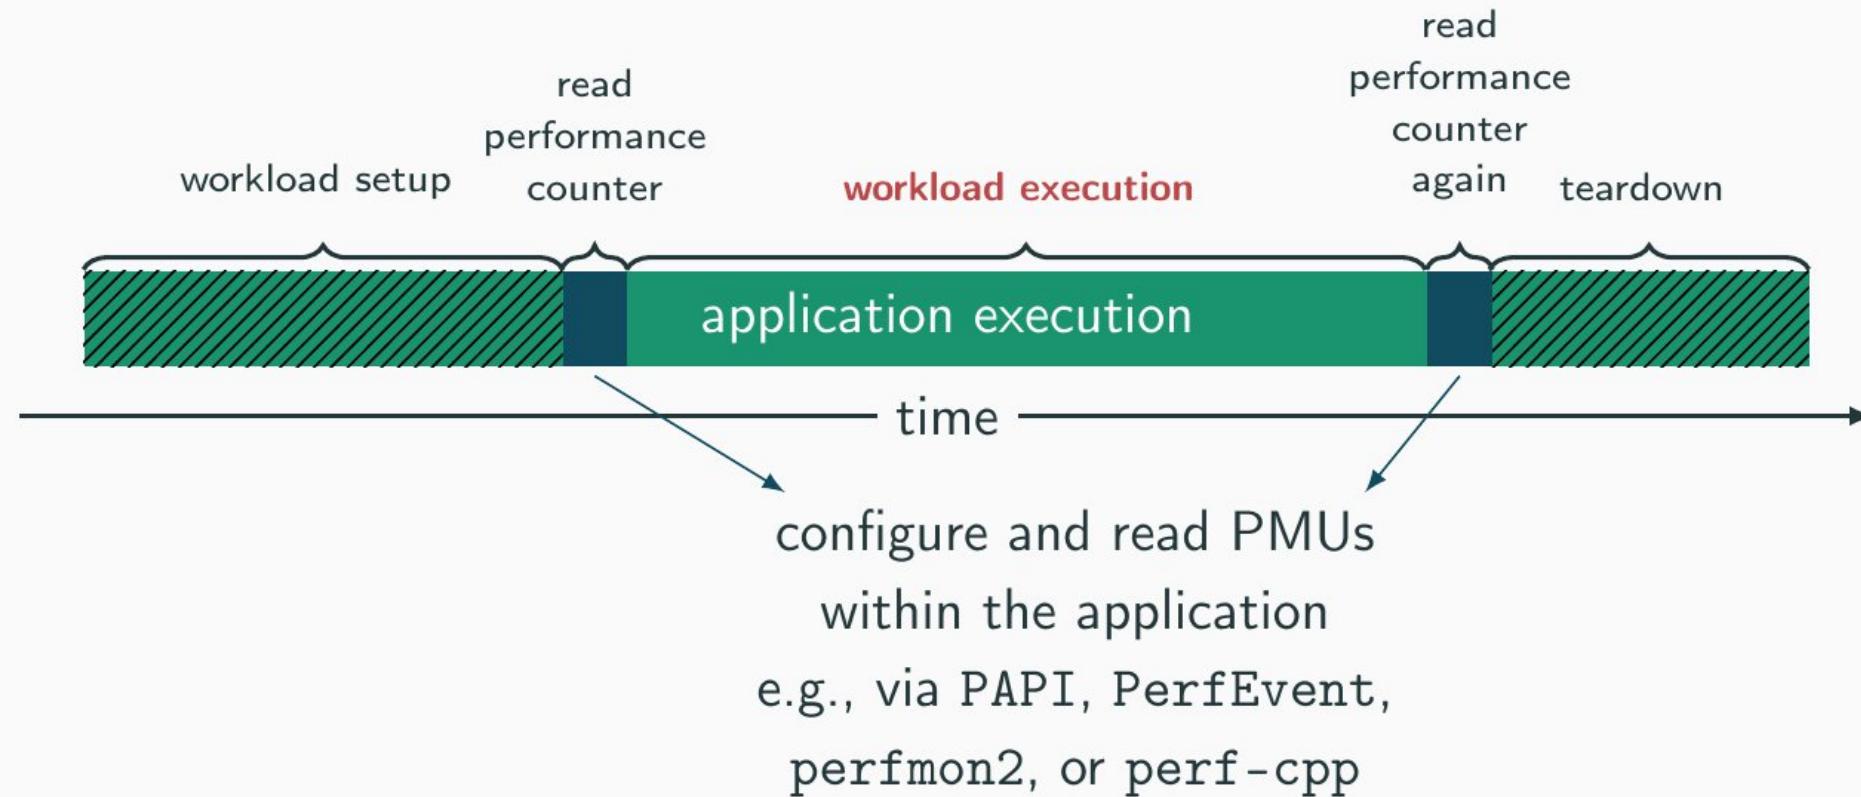
170B	cycles	
111B	instructions	# 0.66 insn per cycle
23B	branches	
1B	branch-misses	# 8.22% of all branches
62B	L1-dcache-loads	
4B	L1-dcache-load-misses	# 7.40% of all L1-dcache accesses



Takeaway

- :(Performance Monitoring Units are implementation-dependent
- : The *perf subsystem* unifies configuration and access

In-Application Profiling: Statistics



Profiling: Sampling



- **Idea:** Record periodic samples of the current execution state
- Needs to define
 - Period → *when* to record (e.g., every 4,000th *cycle* or *memory load*)
 - Sample → *what* to include (e.g., *instruction pointer*, *memory address*, ...)



instruction pointer

0x58b140374b60
0x58b140374b64
0x58b140374b60
0x58b140374b6c
...



Binary &
Debug Symbols



```
Samples: 130K of event 'cycles:P', Event count (approx.): 130299130376
Overhead Symbol
24.24% [...] BTree<unsigned long, unsigned long, 256ul>::insert(unsigned long, unsigned long)
20.93% [...] BTree<unsigned long, unsigned long, 256ul>::lookup(unsigned long, unsigned long)
10.70% [...] void std::shuffle<-__gnu_cxx::__normal_iterator<NumericTuple*>, std::vector<Numeric...
10.36% [...] void CoroutineRoundRobinExecutor::execute<unsigned long, unsigned long>(BTree<uns...
3.08% [...] BTree<unsigned long, unsigned long, 256ul>::insert(unsigned long, unsigned long)
1.81% [...] BTree<unsigned long, unsigned long, 256ul>::lookup(unsigned long, unsigned long)
1.71% [...] unlink_chunk.isra.0
1.22% [K] 0xfffffffffb001c77
1.18% [...] unsigned long std::uniform_int_distribution<unsigned long>::operator()<std::linea...
0.93% [...] __int_free_merge_chunk
0.98% [...] __memmove_avx2_unaligned_erns
0.08% [...] __int_free_maybe_consolidate
0.86% [...] __int_malloc
0.78% [...] __int_free
0.47% [...] Coroutine::promise_type::operator new(unsigned long)
0.41% [K] 0xfffffffffb038419b
0.37% [...] NumericalWorkloadSet::NumericalWorkloadSet(unsigned long, unsigned long)::$_1::operat...
0.38% [...] __int_malign
0.26% [K] 0xfffffffffb0281248
0.23% [...] BTreeInner<unsigned long, 256ul>::~BTreeInner()
0.19% [...] __memmove_gpt
0.18% [...] __mid_memalign.isra.0
```

PerfEvent

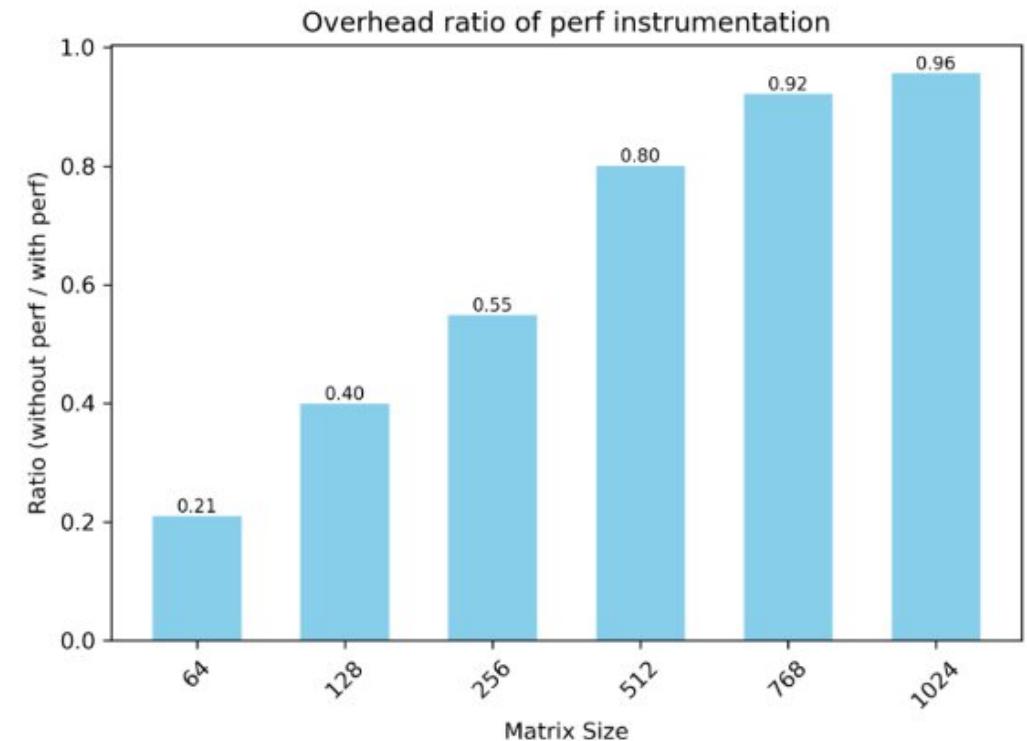
- Great drop-in for simpler projects to weave in perf profiling
- <https://github.com/viktorleis/perfevent>
- A single .hpp header to dump into your library with no complicated dependencies
- We provide it in FOMO/ADMS with the projects already, encourage using it

Likwid

- Another neat Perf-based tool (funded by BMBF btw!)
- <https://github.com/RRZE-HPC/likwid>
- Heavy-duty ‘all-in-one’ – energy consumption, NUMA topology analysis etc
- Can even weave into Fortran code!

Perf overheads

- Likwid overhead shrinks to 5%~10%
(Runtime (without perf) / Runtime (with perf))
when matrix size is larger than 768x768.



Perf-cpp

- <https://github.com/jmuehlig/perf-cpp>
- The project by TU Dortmund
- Includes ‘triggers’: conditions which trigger a PMU capture
- Powerful for more targeted profiling

The memory hierarchy

Practical view

Home cooking is one thing



A factory is a whole other beast



It starts with data

- Much like factories care about supply lines to feed the factory itself, data-intense algorithms need heavy-duty data acquisition mechanisms
- Most modern workloads are **memory-bound** and not compute-bound
- Maximizing memory performance will often speed up the workload

Data structure (DS) properties

5 core metrics

1. Insert Time
2. Update Time
3. Delete Time
4. Access Time
5. Space

Data structure (DS) properties

5 core metrics

1. Insert Time
2. Update Time
3. Delete Time
4. Access Time
5. Space



Severe implications on 1-4
due to the memory
hierarchy

Data structure (DS) properties

5 core metrics

1. Insert Time

2. Update Time

3. Delete Time

4. Access Time

5. Space

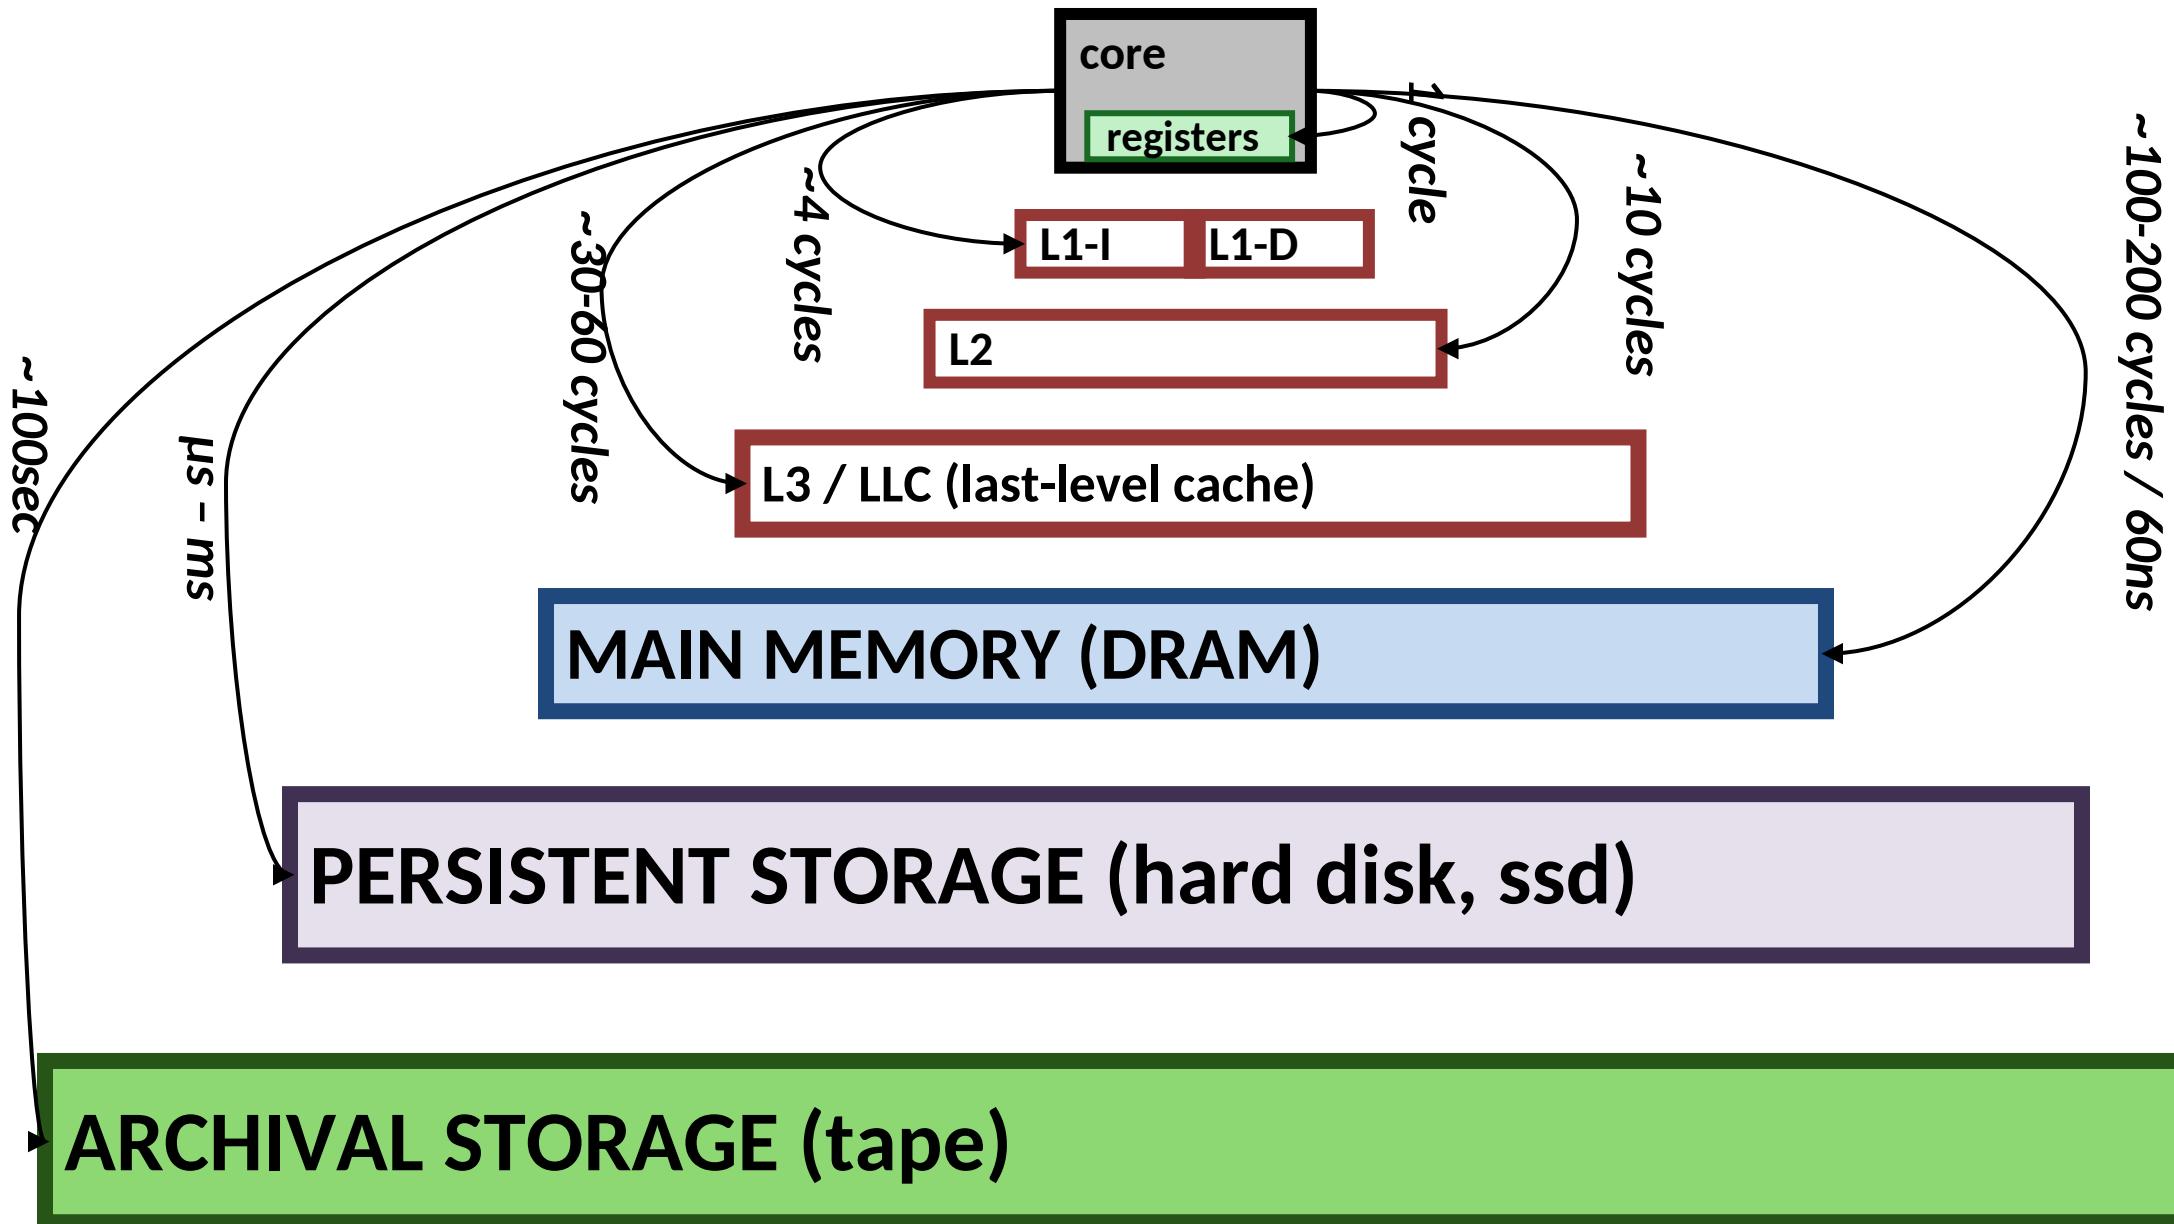
What about these 4 for
sequential access vs
random?

What about worst-case,
average-case, best-
case?

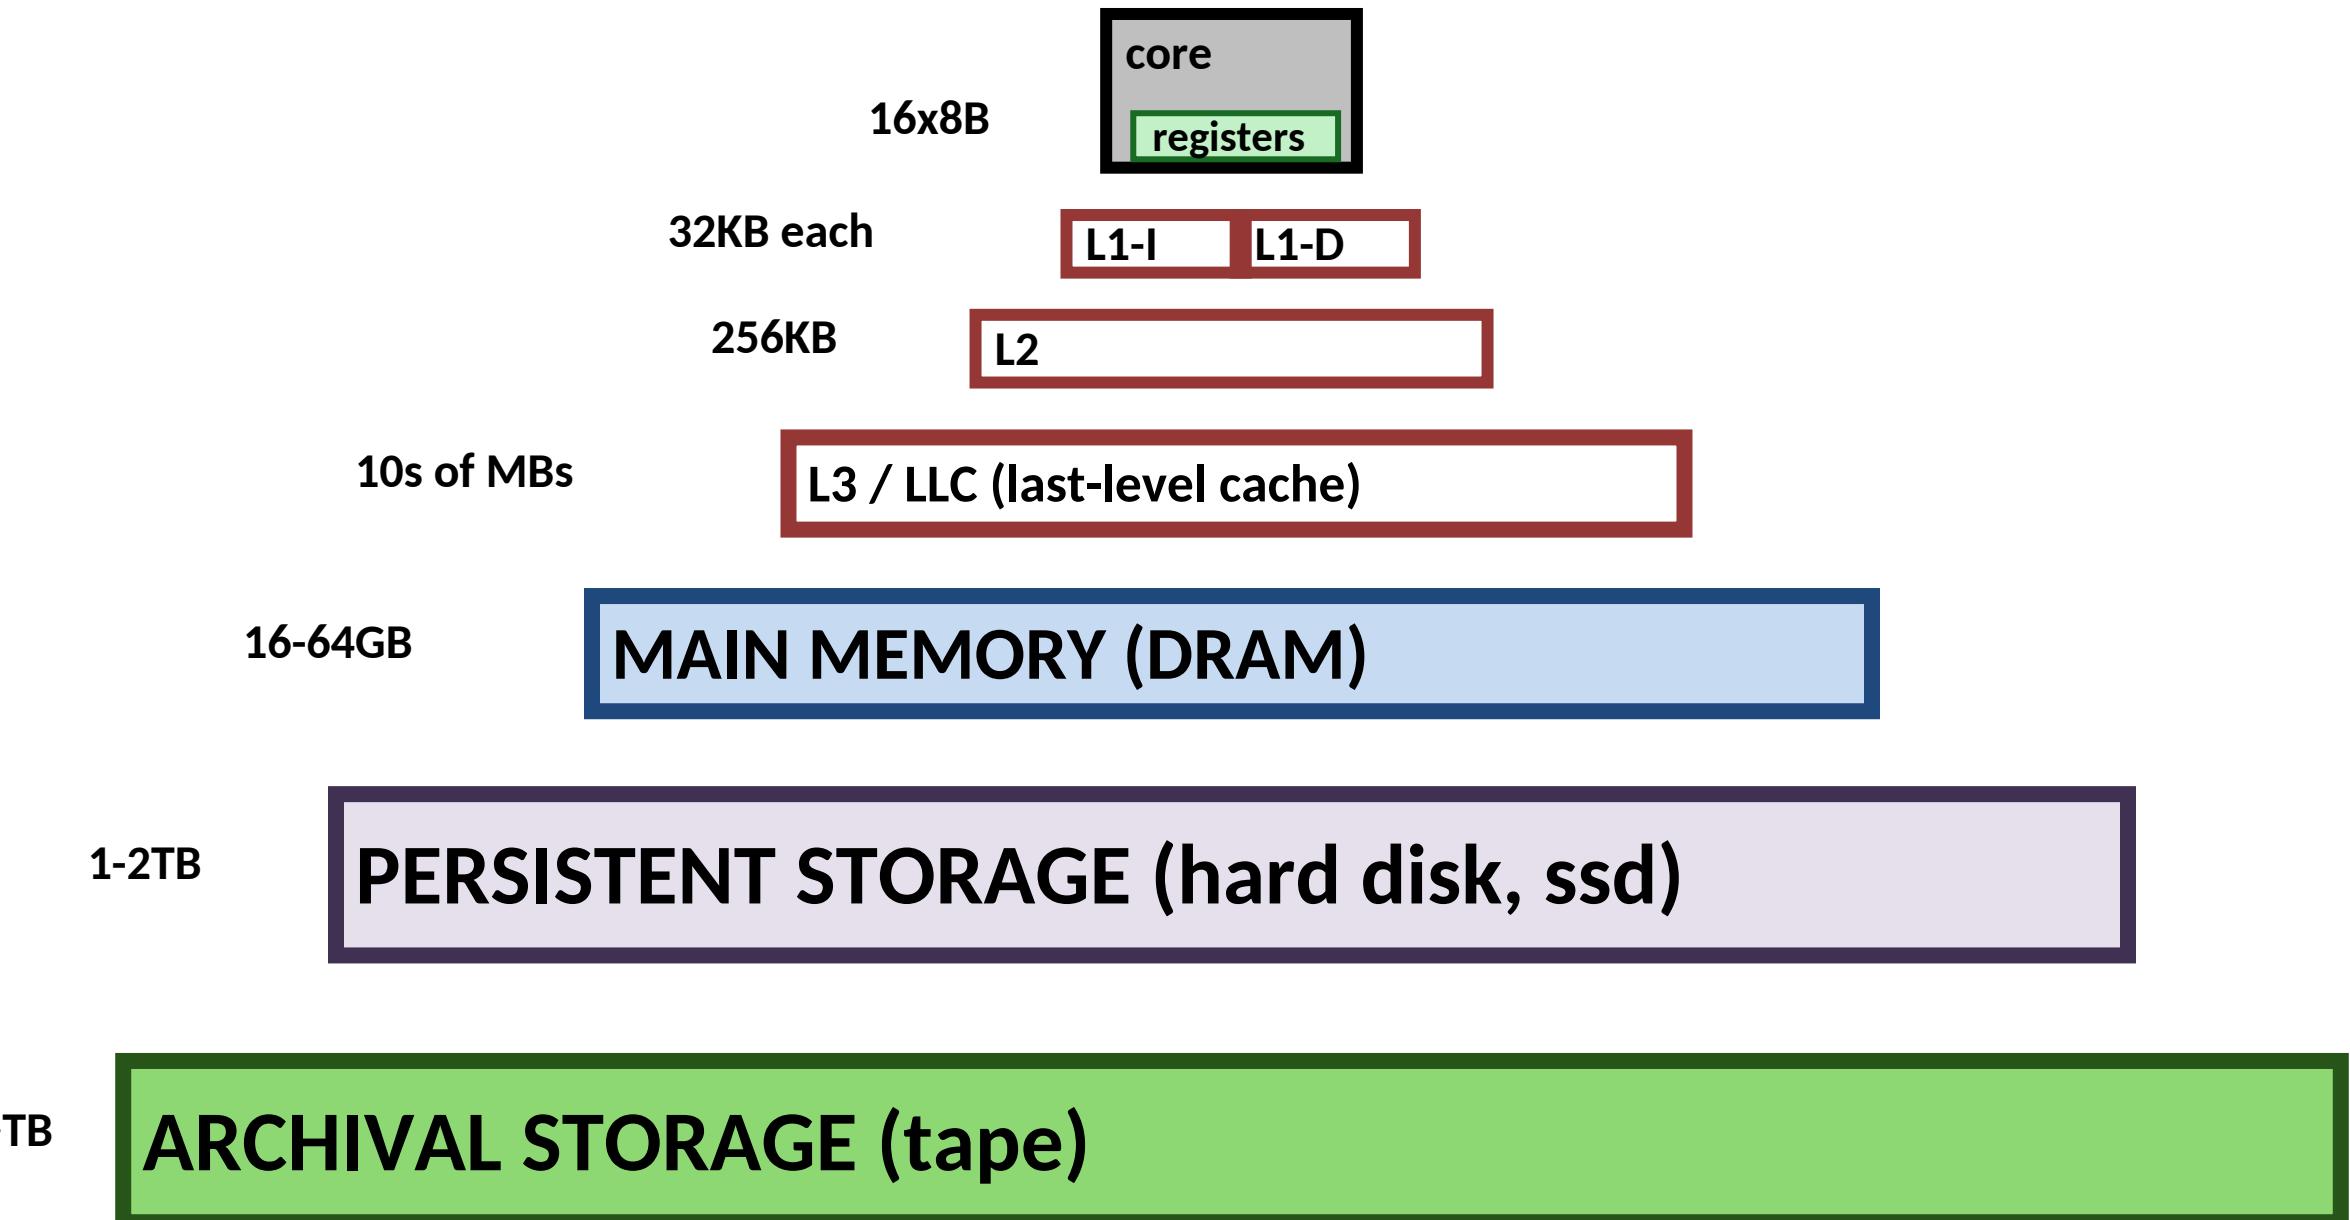
Severe implications on 1-4
due to the memory
hierarchy



(typical) storage hierarchy – access latency



(typical) storage hierarchy – capacity



Why do we care?

- Two data structures with $O(1)$ lookup are unlikely to have the same performance in the real world
- Random access behaves differently to sequential access
- access latency to a memory hurts some DS more than others
- We can trade-off space for decreased latency & higher bw
- We can try to maximize caching effects with reuse
- We can try to exploit locality by prefetching values we might need soon

Why do we care?

- Two data structures with $O(1)$ lookup are unlikely to have the same performance in the real world
- Random access behaves differently to sequential access
- access latency to a memory hurts some DS more than others
- We can trade-off space for decreased latency & higher bw
- We can try to maximize caching effects with reuse
- We can try to exploit locality by prefetching values we might need soon

Key topics we cover next week

Array vs Vector

- Array
 - Static pre-allocated aligned piece of memory
 - Guarantees efficient sequential access
- Lists (vectors)
 - dynamically growing piece of memory
 - No guarantees that memory is aligned
 - Potentially garbage performance (caused by fragmentation)
- Great example: malloc'd arrays in C++ versus std::vector

Array vs Vector

- **Without reserve():**
 - `push_back()` may trigger reallocations
 - Reallocations are costly in terms of performance
- **With reserve():**
 - Allocates required memory upfront
 - Eliminates need for reallocations during insertion
- **Benefits:**
 - Improves insertion performance
 - Reduces memory fragmentation

```
vec.reserve(N);
for (size_t i = 0; i < N; ++i) {
    vec.push_back(i);
}
```

Cache Behavior - Access Patterns

- Access Patterns are key:
 - Sequential Access:
 - Data likely resides in cache & is getting prefetched -> faster access
 - Random Access:
 - Each access is a cache-miss -> slower access
- Use perf to measure cache misses:
 - `perf stat -e cache-misses ./main`

Summary

- Data size, memory alignment and access patterns are key points of consideration for high performance DS
- These are controlled **by you** (or your requirements)
- If you understand them well, you can start considering trade-offs with a hardware awareness

Project 1

Rundown of the projects

Points	Project	Starts today!
15	1. Benchmarking & Memory Hierarchy <ul style="list-style-type: none">Benchmark caches -> Benchmark Data Structures -> Report$O(1) \neq O(1)$	
15	2. Vectorization & Parallelism <ul style="list-style-type: none">Vectorized (SIMD) and Parallelized (multi-threading) implementation of an analytical algorithm<i>Multi-threading isn't always a good idea</i>	
20	3. Accelerators <ul style="list-style-type: none">Optimize a given application for both CPU and GPU and offload when appropriate<i>Offloading isn't always a good idea</i>	
Total: 50		

Learning Objectives

- Accurately measure and analyze hardware memory performance metrics.
- Evaluate and compare the performance of different data structures.
- Clearly and concisely communicate benchmarking results using visual and written formats.
- Explain how the modern hardware memory hierarchy influences the performance of data structures.

How will we do this?

- Benchmark some data structures!
- Make some cool plots!
- Dig into the weeds of how the CPU memory hierarchy works...

Our Data - the humble 64B Node

```
struct alignas(64) Node {  
    uint64_t key;  
    uint64_t data;  
    Node *next;  
    char padding[64 - 16 - sizeof(Node *)];  
};
```

A 64 byte aligned Node
Each lookup = a cache line
read

Makes lookups *very*
memory intense

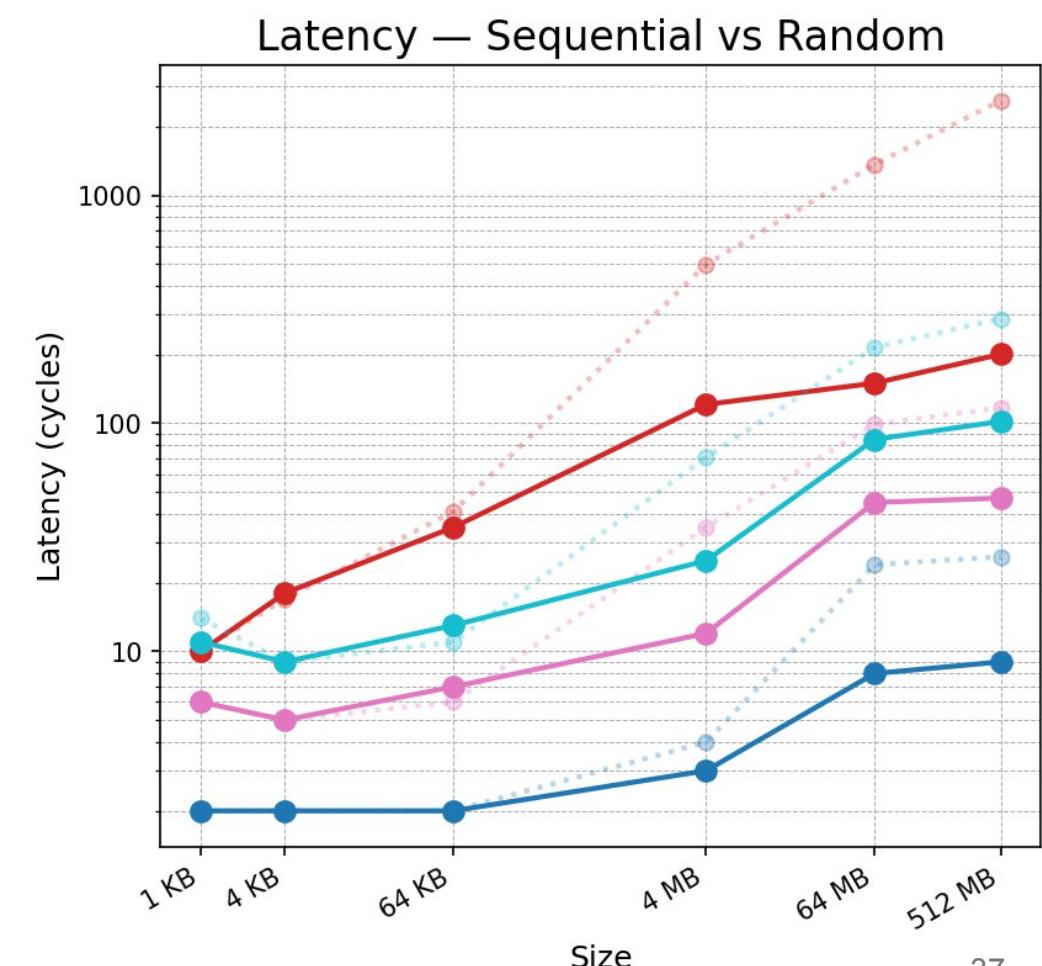
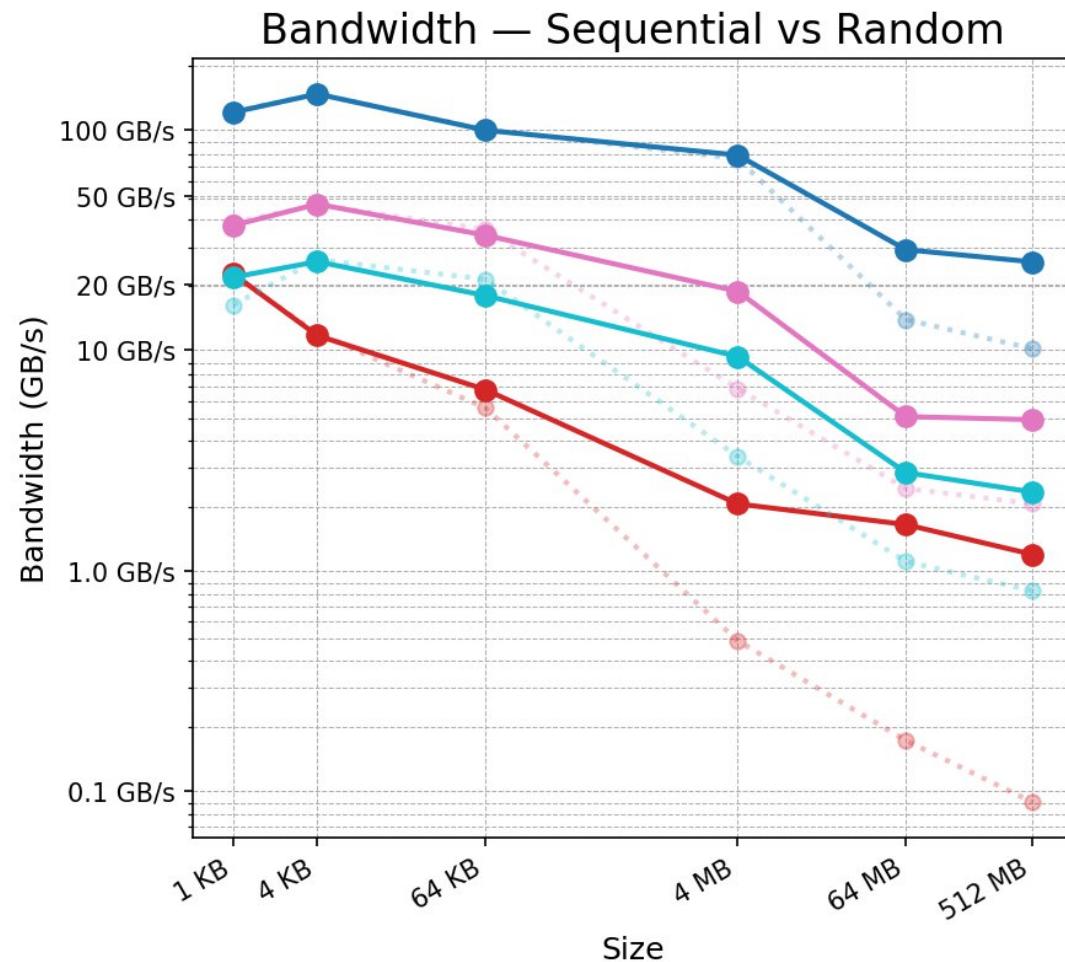
Our contestants

- **Directly accessed array**
 - Designed for dense datasets
 - Pre-initialized memory
 - Keys must be inserted in sequential order
- **Binary search over a list**
 - Designed for sparser datasets (gaps between keys)
 - Still pre-initialized and packed together as a vector for simplicity sake
 - Keys must be inserted in sequential order
- **Chained hash table** (implemented using a vector of vectors)
 - Dynamic size
 - Universal data structure
 - Various trade-offs in terms of chain depth (bin size) and total size of the table
 - Keys could be inserted in random order

Our metrics

- Lookup Bandwidth (average bytes retrieved per second)
- Lookup Latency (average cycles/lookup)

Your Goal



Your Goal - 3 Tasks

1. (3 pts) Evaluate the bandwidth and latency of your memory system
2. (6 pts) Evaluate the bandwidth and latency of 4 distinct data structures
 - Directly accessed array
 - Binary Search over a list
 - Chained Hash with a bin size of 1
 - Chained Hash with a bin size of 16
3. (6 pts) Report: explain how you obtained your results and what do they mean

Your benchmark

- Dataset Generation
- Data Structure Initialization
- Performance measurement (lookups)
- Consider some questions:
 - How do I sufficiently ‘stress’ the data structures to really measure memory costs and not compute costs?
 - How to ‘fairly’ simulate random access?

Some restrictions on key order for your sanity

- All data structures should have N Nodes inserted to them in ascending order by key (even the hash table)
- A key should not be looked up more often than once every N lookups
- All keys should be looked up the same amount of times at the end of the test.

A note on the chained hash table

We ask to evaluate 4 data structures but provide you with 3 classes. What gives?

The **ChainedHashTable** has a hyper parameter – bin size

Use this to initialize two variants of the data structure:

- Chained Hash with a bin size of 1
 - Is this still a hash table?
 - How will it perform relative to an array?
- Chained Hash with a bin size of 16
 - How does it compare to bin size 1?

Final outputs

1. DirectAccessArray
2. BinarySearchArray
3. ChainedHashTable(bin_size=1)
4. ChainedHashTable(bin_size=16)

return {bw_1, bw_2, bw_3, bw_4, lat_1, lat_2, lat_3, lat_4}

Tests

- 2 Basic tests (0.5 points each)
- 8 Advanced tests (1 point each)

The Report

1-2 pages, font size 11

1 point - **explanation** of the benchmarking setup (how the dataset is prepared, how the data structures are evaluated, etc)

1 point - **a plot** showing the bandwidth of all the data structures across varied access patterns and dataset sizes.

1 point - **a plot** showing the latency of all the data structures across varied access patterns and dataset sizes.

1 point - **Compare** the performance of the **DirectAccessArray** and the **ChainedHashTable** with bin size set to 1. Is the performance what you expected? If not, why not.

1 point - **Compare** the performance of the **BinarySearch** with random access versus sequential access. Is the performance what you expected? If not, why not.

1 point - **Compare** the performance of the **ChainedHashTable** with a bin size of 1 versus a bin size of 16. Is the performance what you expected? If not, why not.

Further points

- Full details are in the project description found under this week's moodle section
- We will talk further about the data structures and specifically the effects of caches & memory alignment in closer detail next week
- FORK the project to your own namespace
- The report should be appended to your fork

Hands-on session

Project 1 code walk-through

QA